

Political Bias and Temporal Dynamics in Large Language Models

Ran Lu¹, Hasti Hosseinpour¹, Abderrahmane Moujar¹, Oloruntobi Olutola¹, Jay Saumil Gupte¹

¹Université Paris-Saclay

Abstract

Large Language Models (LLMs) are increasingly deployed as intermediaries for accessing, interpreting, and evaluating political information. Prior work has demonstrated that LLMs exhibit systematic political biases, reflected in ideological alignment, sentiment toward political actors, and framing choices. However, most existing studies treat political bias as a static property, evaluating models at a single point in time. This assumption overlooks the rapid evolution of LLMs, which undergo frequent updates to training data, safety mechanisms, and alignment procedures.

In this work, we investigate political bias as a temporal phenomenon, analyzing how political attitudes expressed by LLMs evolve across model generations. Focusing on longitudinal comparisons within model families, we examine how changes in safety data distributions and alignment interventions correspond to shifts in political stance and sentiment. Our study integrates content-based and sentiment-based bias metrics to provide a dynamic perspective on value alignment in large language models.

1 Introduction

Large Language Models (LLMs) have become central intermediaries in how individuals access, interpret, and evaluate political information. Their widespread use in tasks such as summarization, explanation, sentiment analysis, and question answering grants them substantial influence over contemporary digital public spheres. As these systems are increasingly embedded in search engines, social media platforms, and decision-support tools, their latent political orientations carry meaningful societal and political consequences.

A growing body of evidence suggests that LLMs are not ideologically neutral. Instead, political tendencies emerge from architectural decisions, training data composition, and post-training alignment

strategies (Santurkar et al., 2023; Buyl et al., 2026). When models systematically favor or disfavor particular political ideologies, policies, or actors, they may function as implicit instruments of digital soft power, especially when deployed at scale across languages and geopolitical contexts.

Empirical research further shows that political bias in LLMs is not limited to isolated prompts but constitutes a structural property of model behavior. Studies have identified consistent ideological patterns across economic and social dimensions, as well as systematic sentiment differences toward political entities (Hartmann et al., 2023; Bang et al., 2024). Moreover, LLMs may exhibit bias amplification, whereby biases present in training data are reproduced or intensified in generated outputs, reinforcing existing societal asymmetries (Feng et al., 2023; Zhu et al., 2024).

Despite this progress, most prior work evaluates political bias under a static assumption, treating models as fixed artifacts. This perspective is increasingly inadequate given the rapid pace of LLM development. Model families such as LLaMA undergo substantial changes across generations, including modifications to training corpora, safety data distributions, and reinforcement learning-based alignment procedures (Touvron et al., 2023; Dubey et al., 2024). These interventions plausibly affect not only factual reliability and safety but also political framing and sentiment.

1.1 Research Questions and Contributions

To address this gap, we study political bias through a temporal and longitudinal lens. Specifically, this paper addresses the following research questions:

RQ1: How does political bias in LLM outputs change across successive model generations?

RQ2: To what extent do safety alignment and red-teaming interventions correspond to shifts in political stance and sentiment?

RQ3: Are temporal changes in political bias

uniform across ideological dimensions and political entities, or do they exhibit asymmetric patterns?

Our contributions are threefold:

- We conceptualize political bias in LLMs as a dynamic property, rather than a static characteristic.
- We provide a longitudinal analysis of political bias across model generations using complementary stance- and sentiment-based metrics.
- We empirically link observed bias shifts to documented alignment and safety interventions, contributing to a deeper understanding of value alignment over time.

2 Related Work

Research on political bias in large language models spans ideological measurement, sentiment analysis, temporal evaluation, and mitigation. We review the most relevant work along three axes.

2.1 Measuring Political Bias in Language Models

Early studies mapped LLM outputs onto established ideological frameworks, demonstrating that conversational models express coherent political orientations. Some argue that LLMs reflect the aggregate opinions embedded in their training data (Santurkar et al., 2023), while others situate ChatGPT along a left-libertarian axis using political compass evaluations (Hartmann et al., 2023).

Subsequent work refined bias measurement by distinguishing between content and style. A two-pronged framework analyzing both explicit stance (“what is said”) and framing or lexical polarity (“how it is said”) reveals subtle partisan effects even in ostensibly neutral responses (Bang et al., 2024). Target-oriented sentiment classification further operationalizes political bias by measuring sentiment toward specific political actors or parties (Wang et al., 2025; Chen et al., 2024).

2.2 Temporal Dynamics and Model Evolution

Only recently has attention shifted toward the temporal dimension of LLM behavior. Research shows that LLMs exhibit temporal generalization failures, with performance and biases drifting as political and informational contexts evolve (Zhu et al., 2024). Models implicitly encode the ideological

preferences of their creators, suggesting that updates to training and alignment pipelines may systematically reshape political outputs (Buyl et al., 2026).

Technical reports on the LLaMA family provide concrete evidence of such evolution. Between LLaMA 1, 2, and 3, Meta introduced substantial changes to safety data distributions, reinforcement learning objectives, and red-teaming protocols (Touvron et al., 2023; Dubey et al., 2024). Comparative evaluations indicate that newer models handle politically nuanced and ambiguous content differently, particularly in borderline factual or normative cases (Crum et al., 2024).

2.3 Mitigation and Alignment Interventions

A parallel line of work investigates methods for mitigating political bias. Reinforced calibration techniques have been proposed to reduce ideological skew (Liu et al., 2021), while other research demonstrates that political stance is encoded in internal model activations and can be selectively steered (Banko et al., 2025). These findings suggest that political bias is not merely emergent but amenable to targeted intervention.

However, mitigation introduces trade-offs. Alignment and safety guardrails may reduce expressive or argumentative capacity, potentially altering political nuance (Bonaldi et al., 2024). The LLaMA 3 report explicitly frames safety alignment and red-teaming as central design goals, raising questions about how such interventions affect political expression over time (Dubey et al., 2024).

3 Methodology

3.1 Measuring Political Bias of LLMs

To evaluate the evolution of political bias, we adapt the framework proposed by Elbounani et al. (2025), which quantifies bias through target-oriented sentiment classification (TSC). While their work focuses on a static snapshot of models, we extend this methodology to a temporal dimension to analyze how these biases shift over time.

3.1.1 Data Selection

Entity Selection: We construct a set of 250 political figures sampled from major geopolitical regions to ensure broad ideological and demographic coverage. The selected entities span the full political spectrum from far-left to far-right and are diversified across individual attributes, including gender, ethnicity, socio-economic background, and age.

This diversity is intended to reduce confounding effects caused by over-representation of specific political or demographic groups and to enable a fine-grained analysis of entity-related bias.

Sentence Selection: We select 60 sentences from news-style political discourse, categorized into positive, negative, and neutral sentiments. Consistent with the criteria in Elbouanani et al. (2025), we exclude sentences containing explicit role-specific references (e.g., “the current U.S. president”) to ensure that the political figure’s name is the only varying component, thus maintaining counterfactual consistency during entity substitution.

Examples:

- *Positive:* X contributed to easing tensions during the negotiations.
- *Negative:* X was criticized for mishandling the crisis.
- *Neutral:* X addressed the issue during a public speech.

Multilingual Adaptation: To examine whether political bias is language-dependent, we conduct experiments in multiple languages. For instance, U.S. political figures are evaluated using both English and Chinese prompts, allowing us to assess whether sentiment predictions and bias patterns remain consistent across linguistic contexts. Furthermore, recognizing that training sets in different languages may harbor specific biases—for instance, evaluations of U.S. political figures in English versus Chinese may reflect biases inherent to language-specific corpora rather than the model architecture itself—we explicitly incorporate this factor into our data collection and cross-linguistic comparisons.

3.1.2 Metrics Explanation

We operationalize political bias as prediction inconsistency (IC), a metric defined by Elbouanani et al. (2025) to measure the instability of sentiment predictions when target entities are swapped. For a given sentence, a politically unbiased model is expected to produce identical sentiment predictions regardless of the political entity being substituted. Deviations from this behavior indicate entity-related bias.

We quantify this phenomenon using IC, defined as the average entropy of sentiment predictions across entities for each sentence. Higher entropy

values correspond to greater instability in predictions and, consequently, stronger political bias.

To calculate IC, we first perform entity replacement for each neutral sentence to generate a prediction set. We then derive the probability distribution for each sentiment label $l \in \{positive, neutral, negative\}$ using:

$$P(l | s) = \frac{|\{e \in E : pred(s, e) = l\}|}{|E|} \quad (1)$$

where E is the set of all entities, s is the sentence template, and $pred(s, e)$ returns the model’s predicted sentiment label for sentence s with entity e .

These probabilities are used to compute the Shannon entropy for a specific sentence via the formula:

$$H(s) = - \sum_{l \in L} P(l | s) \log P(l | s) \quad (2)$$

where $L = \{positive, neutral, negative\}$.

Finally, the global bias score is determined by averaging the entropy values across all m sentences in the experiment:

$$IC = \frac{1}{m} \sum_{i=1}^m H(s_i) \quad (3)$$

where an IC value of 0 represents a perfectly unbiased and consistent model.

3.1.3 Experiment Setup

We employ a 9-shot prompting strategy, providing nine labeled sentiment classification examples before each query. This few-shot setting is chosen to stabilize model behavior and reduce variability caused by prompt underspecification. To further constrain the output space, we explicitly instruct the model to produce one of three labels only: positive, neutral, or negative.

Consistent with the recommendations of Röttger et al. (2024), we acknowledge that LLM evaluations of values and opinions are highly sensitive to prompt constraints. To ensure the robustness of our results across the time dimension, we incorporate extensive robustness tests as suggested, making local rather than global claims about the manifested political biases to avoid overgeneralization.

3.2 Measuring Political Bias Over Time

3.2.1 Research Objective and Extension Logic

While Section 3.1 establishes a framework for measuring political bias through Target-Oriented Sentiment Classification (TSC) in a static context, this

section extends the analysis to the temporal dimension. We treat Large Language Models (LLMs) as evolving systems rather than static tools. The primary objective is to determine whether political bias—defined as Prediction Inconsistency (IC)—is being mitigated or amplified as model families transition through successive generations and alignment phases.

To ensure experimental control, this section reuses the 250 entities, 60 neutral sentence templates, and multilingual prompts defined in the previous section. By keeping these constants, we isolate the *Model Version* as the primary independent variable, allowing for a pure measurement of how model evolution affects ideological consistency.

3.2.2 Longitudinal Model Selection and Lineages

To capture a representative trajectory of AI evolution, we categorize models into “Lineages.” This selection process allows us to observe bias across three distinct evolutionary paths:

- **Generational Scaling:** tracking architectures across releases.
- **Alignment Progression:** comparing Base vs. Chat (RLHF-tuned) models within the same family.
- **Geopolitical Diversification:** including Western (Mistral, Llama), Eastern (Qwen), and Middle Eastern (ALLaM) families.

3.2.3 Operationalizing the Evolutionary Metrics

To quantify “Bias Drift” over time, we introduce three longitudinal metrics:

1. Bias Velocity (Slope β): The rate of change in IC across releases:

$$\beta = \frac{IC(V_n) - IC(V_{n-1})}{\Delta t} \quad (4)$$

Interpretation: $\beta < 0 \rightarrow$ Bias Decay; $\beta > 0 \rightarrow$ Bias Accretion.

2. Alignment Delta (Δ_{aln}): The effect of alignment tuning within a generation:

$$\Delta_{aln} = IC(ChatVersion) - IC(BaseVersion) \quad (5)$$

Interpretation: $\Delta_{aln} > 0 \rightarrow$ alignment reduces neutrality; $\Delta_{aln} < 0 \rightarrow$ alignment improves neutrality.

3. Cross-Family Convergence (C): The variance of IC across families at time step s :

$$C(s) = Var(IC(F_1, s), \dots, IC(F_k, s)) \quad (6)$$

Interpretation: $C \downarrow \rightarrow$ convergence toward “Global Neutral” baseline; $C \uparrow \rightarrow$ ideological divergence across regions.

3.2.4 Experimental Procedure and Implementation

The experiment follows a deterministic four-step execution plan:

- 1. Deterministic Sampling:** Temperature = 0 for all queries.
- 2. Snapshot Testing:** Versions tested on identical prompt sets.
- 3. IC Sequence Generation:** $S_F = \{IC_{v_1}, IC_{v_2}, \dots, IC_{v_n}\}.$
- 4. Comparative Attribution:** Sequences grouped by region and alignment.

3.2.5 Statistical Visualization and Interpretation

Results are visualized using Bias Trajectory Maps, where the x-axis represents release date and the y-axis represents IC score.

- **Divergence Analysis:** Trajectories moving upward indicate models grow more opinionated.
- **Stability Analysis:** Flat IC suggests political bias originates from pre-training and resists alignment tuning.

References

- Yejin Bang, Delong Chen, Nayeon Lee, and Pascale Fung. 2024. Measuring political bias in large language models: What is said and how it is said. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*.
- D. Banko and 1 others. 2025. Steering towards fairness: Mitigating political stance bias in LLMs. In *Proceedings of the Annual Meeting of the Association for Computational Linguistics*.
- H. Bonaldi and 1 others. 2024. Alignment and safety guardrails in large language models. In *Proceedings of ACL*.

- M. Buyl and 1 others. 2026. Large language models reflect the ideology of their creators. *Artificial Intelligence*, 2(1).
- Yufei Chen and 1 others. 2024. Unpacking political bias in large language models: A cross-model comparison on U.S. politics. *arXiv preprint arXiv:2412.16746*.
- C. Crum and 1 others. 2024. LLaMA 3 vs. state-of-the-art large language models: Performance in detecting nuanced fake news. *Computers, MDPI*.
- Abhimanyu Dubey and 1 others. 2024. The LLaMA 3 herd of models. *arXiv preprint*.
- Amine Elbouanani, Emma Dufraisse, and Adrian Popescu. 2025. Analyzing political bias in LLMs via target-oriented sentiment classification. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 15476–15505.
- Shangbin Feng and 1 others. 2023. From pretraining data to language models to downstream tasks: Tracking the trails of political biases. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics*.
- Jochen Hartmann and 1 others. 2023. The political ideology of conversational AI. *Royal Society Open Science*.
- Ruibo Liu and 1 others. 2021. Mitigating political bias in language models through reinforced calibration. In *Proceedings of AAAI*.
- Paul Röttger, Valentin Hofmann, Valentina Pyatkin, Musashi Hinck, Hannah Rose Kirk, Hinrich Schütze, and Dirk Hovy. 2024. Political compass or spinning arrow? Towards more meaningful evaluations for values and opinions in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics*.
- Shibani Santurkar and 1 others. 2023. Whose opinions do language models reflect? In *Proceedings of ICML*.
- Hugo Touvron and 1 others. 2023. LLaMA 2: Open foundation and fine-tuned chat models. *Meta AI Research*.
- Yiming Wang and 1 others. 2025. Analyzing political bias in LLMs via target-oriented sentiment classification. In *ACL Anthology*.
- Chen Zhu and 1 others. 2024. Is your LLM outdated? a deep look at temporal generalization. *arXiv preprint arXiv:2405.08460*.