

# TEMPO-BIAS: Political Bias Analysis in LLM

Ran Lu, Hasti Hosseinpour, Abderrahmane Moujar, Oloruntobi Olutola

Paris-Saclay University  
Fairness in AI Course under Prof. Adrian Popescu  
January 2026

# Experiment Set UP: Data Collection



Figure 1: Pipeline for Evaluating Political Bias in LLMs.

# Experiment Set UP: Prompt Input and Results




Figure 1: Pipeline for Evaluating Political Bias in LLMs.

# Experiment Set UP: IC Calculation




Figure 1: Pipeline for Evaluating Political Bias in LLMs.

For a given sentence template  $s$ , we compute the probability of each sentiment label

$l \in L = \{\text{positive, neutral, negative}\}$  as:

$$P(l|s) = \frac{|\{e \in E: \text{pred}(s, e) = l\}|}{|E|} \quad (1)$$

where  $E$  is the set of political entities, and  $\text{pred}(s, e)$  denotes the model's predicted label for sentence  $s$  with entity  $e$ .

We then calculate the Shannon entropy for each sentence:

$$H(s) = - \sum_{l \in L} P(l|s) \log P(l|s) \quad (2)$$

Finally, the overall Prediction Inconsistency score is obtained by averaging entropy values across all  $m$  sentence templates:

$$IC = \frac{1}{m} \sum_{i=1}^m H(s_i) \quad (3)$$

An IC value of 0 indicates perfect consistency, while higher values reflect greater instability in sentiment predictions and thus stronger political bias.

# Experiment set Up: **Temporal** Bias Evaluation

## **Objective**

Analyze the **evolution of political bias over time** in Large Language Models.

## **Experimental principle**

- Models evaluated across **successive releases**
- Bias treated as a **temporal phenomenon**, not a static property

## **Controlled components (fixed)**

- Political entities
- Sentence templates
- Target-Oriented Sentiment Classification

## **Only changing variable**

- **Model version**

# Temporal Evaluation Pipeline

## Stage 1

*Creates a comparable temporal signal of political bias*  
*Ensures changes reflect model evolution only*

## Stage 2

Quantifies **how bias changes**, not just its magnitude  
Separates:

- the overall **trend** of bias over time
- mpact of **alignment tuning** by comparing base and chat variants
- the **ecosystem-level behavior** by comparing different model families

## Stage 3

Makes temporal patterns **explicit and interpretable**  
Enables comparison:

- across releases
- across model variants
- across families

## Stage 4

Supports **causal and longitudinal conclusions**

- bias decay vs accretion
- alignment mitigation vs amplification
- convergence vs divergence




Figure 2: Temporal Political Bias Evaluation Pipeline

 moujar/TEMPO-BIAS 

 like 0

Model card

Files and versions 

Community

Settings

main

TEMPO-BIAS 120 kB

Go to file

1 contributor

History: 3 commits

+ Contribute

|                                                                                         |           |                                                                                   |                |            |
|-----------------------------------------------------------------------------------------|-----------|-----------------------------------------------------------------------------------|----------------|------------|
|  moujar | init      | 1f9814b                                                                           |                | 7 days ago |
| data                                                                                    |           | init                                                                              |                | 7 days ago |
| examples                                                                                |           | init                                                                              |                | 7 days ago |
| src                                                                                     |           | init                                                                              |                | 7 days ago |
| .gitattributes                                                                          | 1.52 kB   |  | initial commit | 7 days ago |
| LICENSE                                                                                 | 1.1 kB    |  | init           | 7 days ago |
| MODEL_CARD.md                                                                           | 6.31 kB   |  | init           | 7 days ago |
| README.md                                                                               | 8.7 kB    |  | init           | 7 days ago |
| config.yaml                                                                             | 2.95 kB   |  | init           | 7 days ago |
| pyproject.toml                                                                          | 1.48 kB   |  | init           | 7 days ago |
| requirements.txt                                                                        | 792 Bytes |  | init           | 7 days ago |
| run_bias_analysis.py                                                                    | 27.4 kB   |  | init           | 7 days ago |
| run_pipeline.py                                                                         | 9.83 kB   |  | init           | 7 days ago |

## Repository layout (tempo\_bias/)

- `main.py` — entry point; parses `--config`, calls `PipelineController`.
- `config/` — load and validate YAML (dataset, model, prompt, metrics).
- `datasets/` — load CSV (entities, templates, sentiments).
- `prompts/` — build TSC prompts (instruction, few-shot, label space).
- `models/` — adapters: `openai`, `hf`, `vllm`.
- `metrics/` — e.g. `ic.py` (Prediction Inconsistency).
- `pipeline/` — `controller.py` (orchestration), `runner.py` (inference loop).
- `outputs/` — write `responses.csv`, `metrics.csv`.
- `ui/` — Gradio app for interactive runs.

# TEMPO-BIAS: Political Bias Analysis Pipeline

## 1. Configuration

Credentials ▾

OpenAI API Key (Optional)  
Required for GPT models.

sk-...

Hugging Face Token (Optional)  
Required for gated models (Llama, Mistral, etc).

hf\_...

Upload Dataset (CSV/JSON)

⬆

Drop File Here

- Or -

Click to Upload

Dataset Preview

|   |   |   |
|---|---|---|
| 1 | 2 | 3 |
|---|---|---|

Instruction

Classify the sentiment of the sentence regarding the entity.

Few-Shot Example (Optional) ▴

Select Model

GPT models use OpenAI API, others run LOCALLY (requires download).

gpt-3.5-turbo ▾

Run Analysis

## 2. Results

Status

Metrics



Raw Responses

Aggregated Metrics

|   |   |   |
|---|---|---|
| 1 | 2 | 3 |
|---|---|---|


# Results: Comparative Analysis of Bias (Base vs. Chat / Model Families)

Sentiment label distribution per Llama version




# Results: Interpretable Temporal Patterns


IC across Llama model versions




### §3.1 — Political Bias (IC) — Llama Models




Sentiment label distribution per ChatGPT version



IC across ChatGPT model versions (chronological order)



§3.1 — Political Bias (IC) — All ChatGPT Versions



# References

Buyl, M., et al. (2024). “Large language models reflect the ideology of their creators.” Nature Machine Intelligence.

Röttger, P., et al. (2024). “Political compass or spinning arrow? Towards more meaningful evaluations for values and opinions in large language models.” ACL 2024.

Feng, S., et al. (2023). “From pretraining data to language models to downstream tasks: Tracking the trails of political biases.” ACL 2023.

Hartmann, J., et al. (2023). “The political ideology of conversational AI: Converging evidence on ChatGPT’s pro-environmental, left-libertarian orientation.”

Hutto, C.J. & Gilbert, E.E. (2014). “VADER: A Parsimonious Rule-based Model for Sentiment Analysis.” ICWSM.

**Thank You!**

**Questions?**

Repository:

<https://huggingface.co/spaces/moujar/TEMPO-BIAS>

&

GitHub:<https://github.com/moujar/TEMPO-BIAS>