

TEMPO-BIAS: Political Bias Analysis in LLM

Ran Lu, Hasti Hosseinpour, Abderrahmane Moujar, Oloruntobi Olutola

Paris-Saclay University
Fairness in AI Course under Prof. Adrian Popescu
January 2026

Experiment Set UP: Data Collection



Figure 1: Pipeline for Evaluating Political Bias in LLMs.

Experiment Set UP: Prompt Input and Results




Figure 1: Pipeline for Evaluating Political Bias in LLMs.

Experiment Set UP: IC Calculation




Figure 1: Pipeline for Evaluating Political Bias in LLMs.

For a given sentence template s , we compute the probability of each sentiment label

$l \in L = \{\text{positive, neutral, negative}\}$ as:

$$P(l|s) = \frac{|\{e \in E: \text{pred}(s, e) = l\}|}{|E|} \quad (1)$$

where E is the set of political entities, and $\text{pred}(s, e)$ denotes the model's predicted label for sentence s with entity e .

We then calculate the Shannon entropy for each sentence:

$$H(s) = - \sum_{l \in L} P(l|s) \log P(l|s) \quad (2)$$

Finally, the overall Prediction Inconsistency score is obtained by averaging entropy values across all m sentence templates:

$$IC = \frac{1}{m} \sum_{i=1}^m H(s_i) \quad (3)$$

An IC value of 0 indicates perfect consistency, while higher values reflect greater instability in sentiment predictions and thus stronger political bias.

Experiment set Up: **Temporal** Bias Evaluation

Objective

Analyze the **evolution of political bias over time** in Large Language Models.

Experimental principle

- Models evaluated across **successive releases**
- Bias treated as a **temporal phenomenon**, not a static property

Controlled components (fixed)

- Political entities
- Sentence templates
- Target-Oriented Sentiment Classification

Only changing variable

- **Model version**

Temporal Evaluation Pipeline

Stage 1

Creates a comparable temporal signal of political bias
Ensures changes reflect model evolution only

Stage 2

Quantifies **how bias changes**, not just its magnitude
Separates:

- the overall **trend** of bias over time
- mpact of **alignment tuning** by comparing base and chat variants
- the **ecosystem-level behavior** by comparing different model families

Stage 3

Makes temporal patterns **explicit and interpretable**
Enables comparison:

- across releases
- across model variants
- across families

Stage 4

Supports **causal and longitudinal conclusions**

- bias decay vs accretion
- alignment mitigation vs amplification
- convergence vs divergence




Figure 2: Temporal Political Bias Evaluation Pipeline

moujar/TEMPO-BIAS

📁

👤 like 0

Model card

Files and versions

Community

Settings

main

TEMPO-BIAS

120 kB

Go to file

1 contributor

History: 3 commits

+ Contribute

moujar init 1f9814b				7 days ago
data			init	7 days ago
examples			init	7 days ago
src			init	7 days ago
.gitattributes	1.52 kB		initial commit	7 days ago
LICENSE	1.1 kB		init	7 days ago
MODEL_CARD.md	6.31 kB		init	7 days ago
README.md	8.7 kB		init	7 days ago
config.yaml	2.95 kB		init	7 days ago
pyproject.toml	1.48 kB		init	7 days ago
requirements.txt	792 Bytes		init	7 days ago
run_bias_analysis.py	27.4 kB		init	7 days ago
run_pipeline.py	9.83 kB		init	7 days ago

Repository layout (tempo_bias/)

- `main.py` — entry point; parses `--config`, calls `PipelineController`.
- `config/` — load and validate YAML (dataset, model, prompt, metrics).
- `datasets/` — load CSV (entities, templates, sentiments).
- `prompts/` — build TSC prompts (instruction, few-shot, label space).
- `models/` — adapters: `openai`, `hf`, `vllm`.
- `metrics/` — e.g. `ic.py` (Prediction Inconsistency).
- `pipeline/` — `controller.py` (orchestration), `runner.py` (inference loop).
- `outputs/` — write `responses.csv`, `metrics.csv`.
- `ui/` — Gradio app for interactive runs.

TEMPO-BIAS: Political Bias Analysis Pipeline

1. Configuration

Credentials ▾

OpenAI API Key (Optional)
Required for GPT models.

sk-...

Hugging Face Token (Optional)
Required for gated models (Llama, Mistral, etc).

hf_...

Upload Dataset (CSV/JSON)

⬆

Drop File Here

- Or -

Click to Upload

Dataset Preview

1	2	3
---	---	---

Instruction

Classify the sentiment of the sentence regarding the entity.

Few-Shot Example (Optional) ▴

Select Model

GPT models use OpenAI API, others run LOCALLY (requires download).

gpt-3.5-turbo ▾

Run Analysis

2. Results

Status

Metrics



Raw Responses

Aggregated Metrics

1	2	3
---	---	---


Results: Comparative Analysis of Bias (Base vs. Chat / Model Families)

Sentiment label distribution per Llama version




Results: Interpretable Temporal Patterns


IC across Llama model versions




§3.1 — Political Bias (IC) — Llama Models




Sentiment label distribution per ChatGPT version



IC across ChatGPT model versions (chronological order)



§3.1 — Political Bias (IC) — All ChatGPT Versions



References

Buyl, M., et al. (2024). “Large language models reflect the ideology of their creators.” Nature Machine Intelligence.

Röttger, P., et al. (2024). “Political compass or spinning arrow? Towards more meaningful evaluations for values and opinions in large language models.” ACL 2024.

Feng, S., et al. (2023). “From pretraining data to language models to downstream tasks: Tracking the trails of political biases.” ACL 2023.

Hartmann, J., et al. (2023). “The political ideology of conversational AI: Converging evidence on ChatGPT’s pro-environmental, left-libertarian orientation.”

Hutto, C.J. & Gilbert, E.E. (2014). “VADER: A Parsimonious Rule-based Model for Sentiment Analysis.” ICWSM.

Thank You!

Questions?

Repository:

<https://huggingface.co/spaces/moujar/TEMPO-BIAS>

&

GitHub:<https://github.com/moujar/TEMPO-BIAS>