

Political Bias and Temporal Dynamics in Large Language models

Ran Lu¹, Hasti Hosseinpour¹, Abderrahmane Moujar¹,
Oloruntobi Olutola¹, Jay Saumil Gupte¹

¹Université Paris-Saclay

Abstract

Large Language Models (LLMs) have become central intermediaries in how individuals access, interpret, and evaluate political information. This survey examines the intersection of political bias and temporal dynamics in LLMs, reviewing recent literature on how models express systematic evaluative patterns when discussing political topics. We discuss key concepts including bias amplification, temporal generalization, nostalgia bias, neophilia bias, and temporal degeneration. We identify a critical research gap: political bias is predominantly studied as a static property, while models evolve over time. This motivates the need to integrate temporal dynamics into the study of political bias in LLMs.

1 Introduction

Large Language Models (LLMs) have become central intermediaries in how individuals access, interpret, and evaluate political information. Their widespread use in tasks such as summarization, explanation, and question-answering grants them substantial power in contemporary digital spheres. However, as Buyl et al. (2026) point out, the behavior of LLMs is far from ideologically neutral; instead, it is shaped by architectural choices, training data composition, and post-training alignment procedures. When models systematically favor or disfavor certain political ideologies or figures, they may function as implicit instruments of digital soft power, particularly when deployed at scale across languages and regions.

Empirical evidence increasingly demonstrates that LLMs express consistent evaluative patterns when discussing political topics, suggesting that political bias is not merely a byproduct of isolated prompts but a structural property of model behavior (Buyl et al., 2026). Moreover, prior work has shown that LLMs may exhibit **Bias Amplification**, whereby biases present in the training data are

not only reproduced but potentially intensified in model outputs, thereby reflecting and exacerbating existing societal prejudices (Zhu et al., 2024). This concern motivates a growing body of work that seeks to measure, interpret, and compare political bias across models.

2 Theoretical Framework

Buyl et al. (2026) define ideology as a systematic worldview concerning social organization, power distribution, and moral evaluation. From this perspective, political bias in LLMs manifests through evaluative discourse—how models morally frame political actors, movements, and doctrines. In operational terms, ideology is quantified as positive or negative preferences toward political labels, and usually measured via sentiment and moral evaluation scores.

While political bias research typically treats models as temporally fixed artifacts, Zhu et al. (2024) introduce a complementary conceptual lens centered on temporal generalization. They define temporal generalization as a model’s ability to reason accurately across past, present, and future contexts. Several key concepts emerge from this framework:

- **Temporal Generalization:** LLMs’ ability to align with the contexts of the past, present, and future.
- **Nostalgia Bias:** Over-reliance on historical data when generating text.
- **Neophilia Bias:** Overemphasis on novelty and future trends.
- **Temporal Degeneration:** Decline in LLM’s performance after it is released.

If LLMs degrade unevenly over time, then political biases themselves may shift, intensify, or invert

across successive model generations. This insight motivates the need to integrate temporal dynamics into the study of political bias.

3 Methodological Landscape

Existing studies adopt diverse evaluation strategies to probe LLMs bias and temporal robustness. Table 1 compares the methodologies of the core works reviewed here.

Buyl et al. (2026) adopt a two-step prompting strategy, first eliciting descriptive text and then asking the model to self-assess its moral stance using a Likert scale. This design enables large-scale cross-model comparison but risks underestimating implicit biases due to reliance on model introspection.

Röttger et al. (2024) critically examine commonly used political evaluation tools, demonstrating that multiple-choice surveys such as the Political Compass Test are highly unstable. Minor paraphrases can produce dramatic shifts in inferred political positions, raising concerns about the validity and robustness of forced-choice evaluations.

Zhu et al. (2024) focus more on temporal robustness rather than ideology. Their FreshBench framework reveals systematic performance decay over time, offering valuable reference for longitudinal analysis.

4 Research Gap

Despite methodological advances, current research exhibits a shared limitation: political bias is predominantly studied as a static property of a single model snapshot. Models are evaluated at one point in time, and differences across architectures or developers are interpreted synchronically. There is limited understanding of whether political biases are stable design features or evolving artifacts shaped by iterative model development.

This gap motivates the need for longitudinal studies that track how political biases evolve across model versions and time periods, integrating insights from both political bias research and temporal generalization frameworks.

Limitations

This survey focuses on a limited set of recent works and does not exhaustively cover all approaches to measuring political bias in LLMs. Additionally, the temporal dynamics framework discussed here

requires further empirical validation through longitudinal studies.

References

- Maarten Buyl and 1 others. 2026. Large language models reflect the ideology of their creators. *Artificial Intelligence*, 2(1):7.
- Paul Röttger and 1 others. 2024. Political compass or spinning arrow? towards more meaningful evaluations for values and opinions in large language models. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*.
- Chenxi Zhu and 1 others. 2024. Is your LLM outdated? a deep look at temporal generalization. *arXiv preprint arXiv:2405.08460*.

Study	Approach	Metrics	Limitations
Buyl et al. (2026)	Two-step prompting with Likert scale self-assessment	Sentiment and moral evaluation scores	May underestimate implicit biases due to reliance on model introspection
Röttger et al. (2024)	Critical examination of political evaluation tools	Multiple-choice survey stability	Shows high instability in forced-choice evaluations
Zhu et al. (2024)	FreshBench framework for temporal analysis	Performance decay metrics	Focuses on temporal robustness rather than ideology

Table 1: Comparison of Methodologies in Related Work